

# An Application of the Implicit Function Theorem

Justin Chiu

Cornell Tech

jtc257@cornell.edu

August 25, 2021

## Abstract

Gradient-based learning forms the foundation of modern machine learning, and automatic differentiation allows ML practitioners to easily compute gradients. While automatic differentiation only costs a constant multiple of the time and space required to evaluate a function, it has its limitations. In particular, when evaluating a function itself requires many computations, the direct application of reverse-mode automatic differentiation may quickly run out of memory. In this report, we review the implicit function theorem (IFT) and its use in reducing the cost of computing gradients in scenarios where function evaluation is expensive, focusing on the application of the IFT to differentiating through the solutions of optimization problems.

## 1 Introduction

Gradient-based learning underpins many of the recent successes in machine learning, particularly advances involving neural networks. The key to the success of gradient-based methods is automatic differentiation (AD), which greatly increases the development speed of machine learning research by allowing practitioners to circumvent the error-prone and time-consuming process of computing gradients manually. AD operates by reducing functions into compositions of atomic operations, for which we have a library of derivatives for, and composing those derivatives via the chain rule. The underlying concept behind AD is that a program's execution trace is a valid and useful representation of a function [Griewank and Walther, 2008], where the execution trace is a history of all the computation performed during the evaluation of a program used to express a function.

Storing the execution trace of a program allows AD systems to easily compute derivatives. However, longer execution traces can quickly consume a large amount of memory. Consider an iterative method, such as gradient descent, whose execution trace takes the form of an unrolled loop: Given an initial point  $\theta = x_0 \in \mathbb{R}^n$ , iterates  $x_1, x_2, \dots, x_{K-1}$  and final point  $x = x_K$  are produced by running the gradient descent update for  $K$  iterations, i.e.  $x_{t+1} = x_t + \alpha \nabla f(x_t)$ . In order to differentiate through this procedure with reverse-mode AD (i.e. compute  $\frac{dx}{d\theta}$ ), we have to store all the  $x_k$  iterates as well the computation used to produce them. Thus, the memory complexity of storing this execution trace scales linearly in the number of iterations  $K$  as well as the dimensionality




Figure 1: Reverse-mode automatic differentiation of an iterative method consumes memory linear in the number of iterations, illustrated by the squiggly line from initial point  $\theta$  to the final point  $x$ . Such methods are not uncommon, found in bilevel optimization problems such as hyperparameter optimization and metalearning. The implicit function theorem (IFT) provides conditions under which there exists an implicit function,  $x^*(\theta) = x$  (dashed line), which we can use to compute the derivative  $\frac{dx_K}{d\theta}$  without storing all intermediate computation.

of the iterates  $x_k$ <sup>1</sup>. For large  $K$ , this can be infeasible as the space complexity is at least  $\Omega(Kn)$ , as this ignores the traces produced by the updates themselves. One method for overcoming the space complexity’s dependence on the number of iterations  $K$  in the above example is to use the implicit function theorem (IFT), letting you throw away intermediate computation (e.g. iterates  $x_1$  through  $x_{K-1}$ ) while still being able to compute the derivative  $\frac{dx}{d\theta}$ . Naively, this comes at the cost of  $O(n^2)$  space and  $O(n^3)$  time, trading off the dependence on  $K$  for a larger dependence on  $n$ .<sup>2</sup> We provide an illustration in Fig. 1.

Large execution traces are not uncommon; optimization problems are often solved with iterative methods, resulting in traces very similar to the example given above. In order to speed up solvers for these problems, one could learn an initialization by differentiating through the execution trace of the iterative method [Finn et al., 2017, Kim et al., 2018, Venkataraman and Amos, 2021]. In this report, we will cover the use of the IFT as a method for dealing with the space complexity of AD in exactly these cases. In particular, we will focus on applying the IFT to differentiating the solutions to optimization problems [Amos and Kolter, 2017, Agrawal et al., 2019]. We will walk through an example by applying the IFT to an optimization problem that is equivalent to softmax, then show how the approach is generalized [Amos and Kolter, 2017].

## 2 Related Work

There are a variety of methods for reducing the space limitations of AD, of which we only mention three: checkpointing, reversible computation, and implicit differentiation.

<sup>1</sup> Forward-mode would get away from the linear dependence on  $K$ , but would require  $n = \dim(\theta)$  sweeps. If all sweeps are performed in parallel, forward-mode AD would take  $\Omega(n^2)$  memory as opposed to  $\Omega(Kn)$  for reverse-mode.

<sup>2</sup> This can be reduced by using approximate vector-inverse Hessian products, as we will show later in this report.

The first method, checkpointing, improves space complexity at the cost of time [Griewank and Walther, 2008]. Rather than storing the full execution trace of a program, checkpointing instead recomputes values when needed. This can result in a slowdown due to recomputation, and also requires careful choosing of which part of the trace to checkpoint and recompute.

A second method is reversible computation [Maclaurin et al., 2015, Gomez et al., 2017], which improves space complexity at the cost of expressivity, but not speed. Reversible computation ensures that a function’s derivative depends only on the output, allowing the input to be discarded during function evaluation. This is typically accomplished by ensuring that the input is easily reconstructed from the output, restricting the expressivity of layers.

A third method uses the IFT, which we focus on. Application of the IFT potentially improves space complexity at the cost of stronger assumptions. The IFT gives conditions under which derivatives can be computed independent of intermediate computation, with the primary condition being the characterization of the output as the solution to a system of equations.

One of the main applications where the space complexity of AD limits its scalability is bilevel optimization. Bilevel optimization problems are optimization problems with another nested inner optimization problem embedded within.

The application we focus on in this report is expressing individual layers of a neural network declaratively as the solution of an optimization problem [Amos and Kolter, 2017, Agrawal et al., 2019, Gould et al., 2019]. This allows models to learn, without heavy manual specification, the constraints of the problem in addition to the parameters of the objective. An example of this is learning to play Sudoku from only input-output pairs [Amos and Kolter, 2017].

Other applications that can be formulated as bilevel optimization problems are hyperparameter optimization, metalearning, and variational inference. Hyperparameter optimization formulates hyperparameter tuning as a bilevel optimization problem, as for each hyperparameter configuration a new model must be trained as the inner loop [Maclaurin et al., 2015, Lorraine et al., 2019b,a, Bertrand et al., 2020]. Derivatives must then be propagated through the inner training loop to the outer hyperparameter loop. Similarly, metalearning learns the parameters of a model such that the model is able to quickly be adapted to a new task via gradient descent [Finn et al., 2017, Rajeswaran et al., 2019]. This is accomplished by differentiating through the learning procedure of each new task. Finally, a variant of variational inference follows a very similar format: semi-amortized variational inference (SAVI) aims to learn a model that is able to provide a good initialization for variational parameters that are subsequently updated iteratively to maximize a lower bound objective [Kim et al., 2018]. This is also accomplished by differentiating through the iterative optimization procedure applied to the variational parameters during inference.

In all the above applications, the inner-loop optimization problem is solved with an iterative method, except in rare, simple cases. The IFT reduces the memory footprint of automatic differentiation, which would otherwise be difficult to scale.

### 3 The Implicit Function Theorem

The implicit function theorem (IFT) has a long history, as well as many applications in a wide variety of fields such as economics and differential geometry. For an overview of the history of the




Figure 2: A circle, defined by the relation  $F(\theta, x) = \theta^2 + x^2 - 1 = 0$ . We view  $\theta$  as a parameter,  $x$  as a solution to  $F(\theta, x) = 0$  given  $\theta$  and the pair  $(\theta, x)$  is a solution point. Our goal is to compute the derivative  $\frac{dx}{d\theta}$ . While we cannot differentiate the relation  $F$  directly as it is not a function, we can compute the derivative at the solution point  $p_1 = (\theta_1, x_1)$  using the local parameterization (or solution mapping)  $f_1(\theta) = \sqrt{1 - x^2}$ . This parameterization holds in a neighbourhood around  $p_1$ , visualized as an arc. We can use the IFT to compute  $\frac{dx}{d\theta} = - \left[ \frac{dF(\theta, x)}{dx} \right]^{-1} \frac{dF(\theta, x)}{d\theta}$ , which will agree with the derivative of the local parameterization  $\frac{df_1(\theta)}{d\theta}$ . We cannot use the same parameterization at  $p_2 = (\theta_2, 0)$  as the derivative is undefined. In general, the IFT is most useful in cases more complicated than the unit circle, where local parameterizations are too complex to easily write down.

90 IFT and some of its classical applications in mathematics and economics, see the book by Krantz  
91 and Parks [2003].

92 Given a relation and some technical conditions that we will go over soon, the IFT is used to  
93 identify the existence of local function parameterizations for particular points, and compute the  
94 Jacobians of the local function parameterizations for points where the conditions for the IFT hold.  
95 As a motivating example, consider the unit circle, governed by the relation  $F(\theta, x) = \theta^2 + x^2 - 1 =$   
96  $0$ , which can be interpreted as a system of equations. As  $F$  fails the vertical line test, we cannot  
97 write  $x$  as a function of  $\theta$  globally. This prevents us from taking derivatives, for example  $\frac{dx}{d\theta}$ .  
98 However, we can use local parameterizations:  $f_1(\theta) = \sqrt{1 - x^2}$  if  $x > 0$  or  $f_2(\theta) = -\sqrt{1 - x^2}$  if  
99  $x < 0$ . Note that the local parameterizations are functions that hold only within a neighbourhood of  
100 a particular solution point  $(\theta, x)$ ; there may be solution points where we simply cannot compute a  
101 particular derivative, e.g.  $\frac{dx}{d\theta}$  when  $x = 0$ . These local parameterizations then allow us compute the  
102 derivative  $\frac{dx}{d\theta}$  at particular solution points  $(\theta, x)$  using the corresponding parameterization. When  
103 the conditions of the IFT do not hold a particular point, we cannot use the IFT to ensure the existence  
104 of a local parameterization and therefore cannot use the IFT to compute derivatives. In general, this  
105 does not mean that the derivative does not exist, only that the IFT does not know what it is. See  
106 Fig. 2 for an illustration. The IFT generalizes this example, and formalizes the conditions under  
107 which there exist continuous local parameterizations for a given relation or system of equations.

108 While the unit circle in this example has very simple local parameterizations, in general local

parameterizations can be more complicated.<sup>3</sup> Additionally, the IFT does not give the form of the local parameterizations; it only guarantees the existence of one around a point and a way to compute its derivative. The local parameterization is left implicit, hence the ‘implicit’ in IFT.

Formally, given a system of equations  $F(\theta, x) = \mathbf{0}_m$ ,<sup>4</sup> and a solution point  $(\theta, x) \in \mathbb{R}^n \times \mathbb{R}^m$ , the IFT gives sufficient conditions under which  $x$  can locally be written as a function of just the parameters  $\theta$  within a neighbourhood of the solution point  $(\theta, x)$ . We saw this in the unit circle example, where the local parameterizations were valid around a particular solution point. We refer to this function  $x^*(\theta) = x$  as a solution mapping,  $x$  a solution, and  $\theta$  as parameters. These conditions are as follows:

1. We have a solution point  $(\theta, x)$  that satisfies the system of equations  $F(\theta, x) = 0$ .
2.  $F$  has at least continuous first derivatives:  $F \in \mathcal{C}^1$ .
3. The Jacobian matrix of  $F$  wrt  $x$  evaluated at the solution point  $(\theta, x)$  is nonsingular:  $\det \frac{dF(\theta, x)}{dx} \neq 0$ .

Assuming these conditions hold for  $F$  at  $(\theta, x)$ , the IFT asserts the existence of the implicit solution mapping  $x^*(\theta) : \mathbb{R}^n \rightarrow \mathbb{R}^m$  (at the solution point  $(\theta, x)$ ), and that the derivative of the solution mapping wrt the parameters is given by

$$\frac{dx^*(\theta)}{d\theta} = - \left[ \frac{dF(\theta, x)}{dx} \right]^{-1} \frac{dF(\theta, x)}{d\theta} \in \mathbb{R}^{m \times n}. \quad (1)$$

As long as we have chosen a solution point  $(\theta, x)$  that satisfies the conditions, we can ignore how we arrived at that solution point and still compute the derivative of the solution mapping with the IFT. It is important to note that the IFT is local and applies to particular solution points. This means that we can pick and choose which solution points we restrict our analysis to.

**Notation: Jacobian Matrix** Given a function  $F : \mathbb{R}^n \times \mathbb{R}^m \rightarrow \mathbb{R}^m$ , we denote the Jacobian matrix evaluated at the point  $(\theta, x) \in \mathbb{R}^n \times \mathbb{R}^m$  as

$$\frac{dF(\theta, x)}{d(\theta, x)} = \begin{bmatrix} \frac{dF(\theta, x)}{d\theta} & \frac{dF(\theta, x)}{dx} \end{bmatrix},$$

where we have the matrix of partial derivatives

$$\frac{dF(\theta, x)}{d\theta} = \begin{bmatrix} \frac{dF_1(\theta, x)}{d\theta_1} & \dots & \frac{dF_1(\theta, x)}{d\theta_n} \\ \vdots & \ddots & \vdots \\ \frac{dF_m(\theta, x)}{d\theta_1} & \dots & \frac{dF_m(\theta, x)}{d\theta_n} \end{bmatrix} \in \mathbb{R}^{m \times n},$$

and similarly for  $\frac{dF(\theta, x)}{dx} \in \mathbb{R}^{m \times m}$ .

<sup>3</sup> Recall that a program trace is a valid representation of a function. One could have a local parameterization that is the trace of a long program, such as one that consists of a series of iterative updates.

<sup>4</sup> We denote vectors and matrices of all 1s and 0s by  $\mathbf{1}_S$  and  $\mathbf{0}_S$ , where  $S$  denotes the shape, i.e.  $\mathbf{0}_m \in \mathbb{R}^m$ .

Given the IFT and necessary notation, we can now proceed to use the IFT to compute derivatives of the output of an iterative method, as depicted in Fig. 1. In particular, we focus on computing derivatives of the solution to a constrained optimization problem. We will use the optimality criteria of the optimization problem to define a system of equations, then apply the IFT to compute the Jacobian of the solution wrt the parameters. This methodology allows us to use the solution to an optimization problem as the output of a layer within a neural network, as done in OptNet [Amos and Kolter, 2017].

## 4 Differentiating through Optimization

As an introductory example, we will replace softmax with an equivalent function defined as the output of an optimization problem, then derive derivatives using the IFT. We will start by reviewing softmax and its expression as an optimization problem. After checking the conditions of the IFT hold, we can then compute derivatives. Since the Jacobian of softmax is known, we can directly verify that the IFT gives the correct answer.

### 4.1 Softmax

Softmax is often used to parameterize categorical distributions within neural networks, such as in attention layers. Given  $n$  items with independent utilities, where  $\theta \in \mathbb{R}^n$  indicate preferences, softmax gives the following distribution over items:  $z_i = \text{softmax}(\theta)_i = \frac{\exp(\theta_i)}{\sum_j \exp(\theta_j)}$ , with  $z \in [0, 1]^n$ . While there is a closed-form equation for both softmax and its Jacobian, we use it as an introduction to the mechanism behind OptNet (and other differentiable optimization layers) [Amos and Kolter, 2017, Agrawal et al., 2019].

The output of softmax is the solution of the following constrained optimization problem [Gao and Pavel, 2018]:

$$\begin{aligned} & \text{maximize} && z^\top \theta + H(z) \\ & \text{subject to} && z^\top \mathbf{1} = 1 \\ & && z_i \geq 0, \forall i, \end{aligned} \tag{2}$$

where  $H(z) = -\sum_i z_i \log z_i$  is the entropy. The first term in the objective,  $z^\top \theta$ , is highest when  $z$  points in the same direction as  $\theta$ . Given the constraint that  $z$  must sum to one and have nonnegative entries, maximizing just the first term subject to those constraints results in a solution that picks out the highest component of  $\theta$ , i.e.  $\text{argmax}(\theta)$ . The addition of the entropy term penalizes solutions that put too much mass on just a few items, resulting in a smoothed version of  $\text{argmax}$ . We will refer to this entropy-regularized version as the softmax problem.<sup>5</sup>

Our goal is to compute the Jacobian of softmax  $\frac{dz}{d\theta} = \frac{d\text{softmax}(\theta)}{d\theta}$  using the IFT and the optimization problem above. While this may seem trivial because softmax has a closed form expression for both the output and Jacobian, it is a worthwhile exercise in applying the IFT. Applying the IFT to optimization problems consists of four steps:

<sup>5</sup> Removing the entropy regularization term results in the  $\text{argmax}$  optimization problem: maximize  $z^\top \theta$ , subject to  $z^\top \mathbf{1}_n = 1$  and  $z \succeq 0$ . This is known as the support function of the set  $\{z \mid z^\top \mathbf{1}_n = 1, z \succeq 0\}$ .

- 157 1. Find a solution to the optimization problem.
- 158 2. Write down a system of equations derived from the optimality conditions.
- 159 3. Check that the conditions of the IFT hold.
- 160 4. Compute the derivative of the implicit solution mapping wrt the parameters.

161 We assume the first step has been done for us, and we have a solution  $z$  to the softmax problem.<sup>6</sup>  
 162 We will then use the IFT to compute gradients of  $z$  wrt the parameters  $\theta$  by following the rest of the  
 163 steps.

## 164 Step 2: The KKT conditions determine the system of equations

165 In order to apply the IFT, we need a system of equations for which our outputs of interest are  
 166 solution points for. The goal of this step is to choose a suitable system of equations. This differs  
 167 from the example of the unit-circle, which was already defined by a system of equations.<sup>7</sup>

168 For solutions of optimization problems, the Karush-Kuhn-Tucker (KKT) conditions are a nat-  
 169 ural choice for defining such a system of equations. Given an optimization problem, the KKT  
 170 conditions determine a system of equations that the solution must satisfy based on the optimality  
 171 criteria [Karush, 1939, Kuhn and Tucker, 1951]. They are stationarity (the gradient should be 0 at  
 172 a local optima) and feasibility (the constraints of the problem should not be violated). For a thor-  
 173 ough introduction to the KKT conditions, see chapter 5 of Boyd and Vandenberghe [2004] or the  
 174 Wikipedia article

We will use the KKT conditions of the softmax problem in Eqn. 2 to determine the (vector-  
 valued) function  $F : \mathbb{R}^n \times \mathbb{R}^m \rightarrow \mathbb{R}^m$  in the IFT. First, we introduce dual variables  $u \in \mathbb{R}, v \in \mathbb{R}^n$   
 and write out the Lagrangian:

$$\mathcal{L}(\theta, z, u, v) = z^\top \theta + H(z) + u(z^\top \mathbf{1}_n - 1) + v^\top z.$$

We therefore have the solution point  $(\theta, z, u, v)$ , with parameters  $\theta$  and solution  $x = (z, u, v)$ . We  
 then have the following necessary conditions for a solution  $(z, u, v)$ , i.e. the KKT conditions:

$$\begin{aligned} \frac{d}{dx} \mathcal{L}(\theta, z, u, v) &= \mathbf{0}_n && \text{(stationarity)} \\ u(z^\top \mathbf{1} - 1) &= 0 && \text{(primal feasibility, equality)} \\ \text{diag}(v)z &= \mathbf{0}_n && \text{(complementary slackness)} \\ z &\succeq \mathbf{0}_n && \text{(primal feasibility, inequality)} \\ v &\succeq \mathbf{0}_n && \text{(dual feasibility)} \end{aligned} \tag{3}$$

---

<sup>6</sup> This is trivial for softmax since we can compute it using the closed form expression. However, in more general optimization problems, we would obtain  $z$  from a solver with a large execution trace.

<sup>7</sup> For a given solution point there may be many suitable systems, all of which would be viable as long as they satisfy the conditions of the IFT that we will check in step 3.

175 As we only need a system of equations with  $2n + 1$  equations to determine the  $2n + 1$  solution  
 176 variables  $x = (z, u, v)$ , we use the first three conditions: stationary, primal feasibility (equality),  
 177 and complementary slackness.

In full, the system of equations  $F(\theta, z, u, v) = 0$  we choose for the softmax problem is

$$\begin{aligned}\theta + -\log(z) - 1 + u\mathbf{1}_n + v &= \mathbf{0}_n \\ u(z^\top \mathbf{1}_n - 1) &= 0 \\ \text{diag}(v)z &= \mathbf{0}_n.\end{aligned}\tag{4}$$

178 Note that the first and third equations are vector-valued.

179 This completes the second step, where we chose a subset of the KKT conditions in order to  
 180 produce a nonlinear system of equations. We can now proceed check the conditions of the IFT,  
 181 which will determine whether  $F$  is locally well-behaved at a particular solution point  $(\theta, z, u, v)$ .

### 182 **Step 3: Check that the IFT conditions hold at the solution point**

183 Recall that the IFT requires the following three conditions, which must be checked on a case-by-case  
 184 basis for particular solution points<sup>8</sup>:

- 185 •  $F(\theta, z, u, v) = 0$ ,
- 186 •  $F$  has at least continuous first derivatives,
- 187 •  $\det \frac{dF(\theta, z, u, v)}{d(z, u, v)} \neq 0$ , or equivalently  $\frac{dF(\theta, z, u, v)}{d(z, u, v)}$  is full rank.

188 Returning to the softmax problem, the first condition holds as we have a solution to the optimization  
 189 problem and  $F$  was chosen using the KKT conditions.<sup>9</sup> The second condition also holds, as  $F$  has  
 190 continuous first derivatives. All that remains is to check the third condition, that the Jacobian matrix  
 191  $\frac{dF(\theta, z, u, v)}{d(z, u, v)}$  (evaluated at the solution point) is non-singular.

The Jacobian matrix  $\frac{dF(\theta, z, u, v)}{d(z, u, v)} \in \mathbb{R}^{2n+1 \times 2n+1}$  is given by

$$\frac{dF}{d(z, u, v)} = \begin{bmatrix} \text{diag}(z)^{-1} & -\mathbf{1}_n & -I_{n \times n} \\ u\mathbf{1}_n^\top & z^\top \mathbf{1}_n - 1 & 0 \\ \text{diag}(v) & 0 & \text{diag}(z) \end{bmatrix}.\tag{5}$$

192 Since a solution must be feasible, we know that  $z^\top \mathbf{1} - 1 = 0$  and  $u > 0$ . However, the upper left  
 193 block,  $\text{diag}(z)^{-1}$ , contains a divide-by-zero term if any component  $z_i = 0$ .<sup>10</sup> To avoid this, we  
 194 consider only strictly positive  $z \succ 0$  for the IFT.<sup>11</sup> Constraining to strict positivity for  $z$ , we can

<sup>8</sup>If these conditions do not hold, the IFT does not apply at this particular combination of  $F$  and solution point. The derivative  $\frac{dx}{d\theta}$  may still be computed via other means if it exists.

<sup>9</sup> Recall that softmax also has a closed form expression.

<sup>10</sup> This term was obtained by differentiating the entropy term of the Lagrangian,  $H(z) = \sum_i z_i \log z_i$ . While we could use the convention  $0 \log 0 = 0$ , this does not fix the divide-by-zero issue with the second derivative, which we see here.

<sup>11</sup> We saw a similar issue in the unit circle example, where the derivative  $\frac{dx}{d\theta}$  was undefined when  $x = 0$  (see Fig. 2).



deduce that the Jacobian of  $F$  is full rank and therefore has nonzero determinant. This shows that the conditions of the IFT hold for the solution points that are feasible, optimal, and have strictly positive  $z$ .<sup>12</sup>

#### Step 4: Compute $\frac{dx}{d\theta}$

Now that we have a set of solution points where conditions of the IFT hold, we can use the IFT to compute  $\frac{dz}{d\theta}$ . Recall that we have the solution  $x = (z, u, v)$ ; we will switch to  $x$  for brevity. The second part of the IFT (Eqn. 1) tells us that we can compute the Jacobian of the solution mapping  $\frac{dx}{d\theta} = \frac{dx^*(\theta)}{d\theta} = - \left[ \frac{dF(\theta, x)}{dx} \right]^{-1} \frac{dF(\theta, x)}{d\theta}$ , then pick out the relevant components.

The second term,  $\frac{dF(\theta, x)}{d\theta}$ , is simple. Since  $\theta$  only appears in the first vector-valued function of  $F$  (see Eqn. 4), we have

$$\frac{dF(\theta, x)}{d\theta} = \begin{bmatrix} I_{n \times n} \\ \mathbf{0}_{(n+1) \times (n+1)} \end{bmatrix}. \quad (6)$$

The large amount of sparsity allows us to skip some computation further down.

Next, we have to invert the Jacobian from Eqn. 5:

$$\left[ \frac{dF(\theta, x)}{dx} \right]^{-1} = \begin{bmatrix} \text{diag}(z)^{-1} & -\mathbf{1}_n & -I_{n \times n} \\ u\mathbf{1}_n^\top & x^\top \mathbf{1}_n - 1 & 0 \\ \text{diag}(v) & 0 & \text{diag}(z) \end{bmatrix}^{-1}. \quad (7)$$

The remainder of this section is compute-intensive; feel free to skip ahead to Sec. 4.2 for a discussion on the limitations of applying the IFT. We use the block-wise inversion formula

$$\begin{bmatrix} A & B \\ C & D \end{bmatrix}^{-1} = \begin{bmatrix} (A - BD^{-1}C)^{-1} & 0 \\ 0 & (D - CA^{-1}B)^{-1} \end{bmatrix} \begin{bmatrix} I & -BD^{-1} \\ -CA^{-1} & I \end{bmatrix},$$

where

$$\begin{aligned} A &= \begin{bmatrix} \text{diag}(z)^{-1} & -\mathbf{1}_n \\ u\mathbf{1}_n^\top & 0 \end{bmatrix} & B &= \begin{bmatrix} -I_{n \times n} \\ 0 \end{bmatrix} \\ C &= [\text{diag}(v) \quad 0] & D &= \text{diag}(z). \end{aligned}$$

However, by complementary slackness, we have  $v = 0$ ,<sup>13</sup> reducing the above to

$$\begin{bmatrix} A & B \\ C & D \end{bmatrix}^{-1} = \begin{bmatrix} A^{-1} & 0 \\ 0 & D^{-1} \end{bmatrix} \begin{bmatrix} I_{(n+1) \times (n+1)} & -BD^{-1} \\ 0 & I_{n \times n} \end{bmatrix}.$$

As we are interested in computing  $\frac{dz}{d\theta}$ , rather than the full derivative  $\frac{dx}{d\theta}$  (recall  $x = (z, u, v)$ ), in addition to the sparsity of  $\frac{dF}{d\theta}$ , we only have to solve for the upper-left  $n \times n$  block of  $A^{-1} \in$

<sup>12</sup> Recall that the IFT applies at particular solution points, meaning we can pick and choose which points to analyze. In the case of softmax, we can perform a separate analysis for solution points on the boundary of the constraint set.

<sup>13</sup> We apply the IFT to solutions where  $z \succ 0$  due to a divide-by-zero issue in  $\frac{d^2 H(z)}{dz_i^2} = \frac{1}{z_i}$ .

$\mathbb{R}^{n+1 \times n+1}$ . To do so, we will repeat the same block-wise inverse computation. Let us denote

$$A = \begin{bmatrix} \text{diag}(z)^{-1} & -\mathbf{1}_n \\ u\mathbf{1}_n^\top & 0 \end{bmatrix} = \begin{bmatrix} E & F \\ G & H \end{bmatrix}.$$

First, we compute the Schur complement of  $A$ ,

$$A/E = H - GE^{-1}F = 0 + u\mathbf{1}_n^\top \text{diag}(z)\mathbf{1}_n = uz^\top \mathbf{1}_n. \quad (8)$$

Since  $z$  is feasible, we have  $A/E = u$  due to the equality constraints ( $z$  must sum to 1 as a probability mass function). Then, we have

$$A^{-1} = \begin{bmatrix} \text{diag}(z)^{-1} & -\mathbf{1}_n \\ u\mathbf{1}_n^\top & 0 \end{bmatrix}^{-1} = \begin{bmatrix} E & F \\ G & H \end{bmatrix}^{-1} = \begin{bmatrix} E^{-1} + E^{-1}F(A/E)^{-1}GE^{-1} & -E^{-1}F(A/E)^{-1} \\ -(A/E)^{-1}GE^{-1} & (A/E)^{-1} \end{bmatrix}. \quad (9)$$

Plugging in,

$$\begin{aligned} A^{-1} &= \begin{bmatrix} \text{diag}(z) - \text{diag}(z)\mathbf{1}_n u^{-1} u\mathbf{1}_n^\top \text{diag}(z) & \text{diag}(z)\mathbf{1}_n u^{-1} \\ -u^{-1} u\mathbf{1}_n^\top \text{diag}(z) & u^{-1} \end{bmatrix} \\ &= \begin{bmatrix} \text{diag}(z) - zz^\top & u^{-1}z \\ -z^\top & u^{-1} \end{bmatrix}. \end{aligned} \quad (10)$$

Pulling out the top-left  $n \times n$  block yields the Jacobian  $\frac{dz}{d\theta} = \text{diag}(z) - zz^\top$ , which agrees with directly differentiating softmax [Martins and Astudillo, 2016]. While we were able to take advantage of sparsity in this computation, such sparsity may not be available in general. Without sparsity, inverting the Jacobian of  $F$ ,  $\frac{dF(\theta, x)}{dx}$ , costs  $O(n^3)$  time and  $O(n^2)$  space compared to the  $O(n)$  time and space for computing softmax and its derivative exactly (as well as with the IFT and sparsity).

This concludes the exercise of using the IFT to differentiate through the softmax problem. While softmax itself has a closed form expression for both its output and derivative, we showed that the IFT can also be used to compute its derivative in a solver-agnostic manner. We started by expressing softmax as an optimization problem (Eqn. 2), then, assuming we had a solution to the softmax problem (regardless of how it was obtained), reverse engineered a system of equations using the KKT conditions. We then restricted our attention to the set of solution points where we could successfully apply the IFT, then applied Eqn. 1 to compute the derivative at a solution point.

We can immediately apply the methodology developed in this section to more general optimization problems, as we will explore in Sec. 5. In those problems closed form solutions may not exist, and iterative solvers must be used where applying reverse-mode AD would take  $\Omega(Kn)$  space with  $K$  steps, while the IFT would naively take  $O(n^2)$  space.

## 4.2 Limitations

In order to compute the derivative  $\frac{dx}{d\theta}$ , we had to invert the Jacobian of  $F$  (the system of equations derived from the KKT conditions), i.e. compute  $\left[\frac{dF(\theta, x)}{dx}\right]^{-1}$ . However, the first part of  $F$  (recall from Eqn. 4), the stationarity condition  $\frac{d}{dx}\mathcal{L} = 0$  already involved the Jacobian of the Lagrangian

$\mathcal{L}$ . In general, this means that in order to apply the IFT to solutions of optimization problems, we must compute the inverse Hessian of the Lagrangian (or at least an inverse Hessian-vector-product). The Hessian is a matrix of size  $O(n^2)$ , and inverting this would take  $O(n^3)$  computation.<sup>14</sup> Thankfully, there are relatively cheap ways of approximating this computation, such as with approximate (inverse) Hessian-vector-product techniques [Rajeswaran et al., 2019, Lorraine et al., 2019a]. Lorraine et al. [2019a] in particular show that the derivative can be approximated in  $O(n)$  space using a truncated Neumann series approximation for the inverse Hessian.

Additionally, comparing the speed of computing derivatives via AD on a large trace versus the IFT must be done on a case-by-case basis. For softmax, let us pretend we solved the softmax problem with an iterative method in  $K$  steps. Computing the derivative from the execution trace would take  $O(Kn)$  time using reverse-mode AD, while the naive IFT would take  $O(n^3)$  time, giving an asymptotic speedup only when  $K > n^2$ . In practice, the constant factors in the runtime will play a large factor in which approach is better, as well as any approximations made to make either approach more efficient<sup>15</sup>.

### 4.3 Extensions

The methods we covered can be extended to variations of argmax problems other than the softmax problem.<sup>16</sup> The softmax problem altered the argmax problem by introducing entropy regularization. Rather than regularizing with entropy, one could instead alter the objective to find the Euclidean projection of the parameters onto the probability simplex, resulting in SparseMax [Martins and Astudillo, 2016]. While the output of softmax variants often have a closed form expression, the IFT provides another way of deriving their Jacobians and could potentially pave the way for differentiating through argmax problems that do not have closed-form expressions.

More generally, the IFT can be applied to cases where, unlike softmax, we do not have an explicit functional form (i.e., the unit circle), and outputs are governed only by a system of equations. This includes more general optimization problems, such as quadratic programs [Amos and Kolter, 2017] or other convex optimization problems [Agrawal et al., 2019].

## 5 OptNet

OptNet generalizes the methodology applied above to the softmax problem by extending the optimization problems considered, in particular including parameterized constraints. This allows us to learn not only the objective, but also the constraints. Explicitly incorporating families of constraints in models with optimization layers to allows them to perform well on tasks with rigid constraints, such as learning to play Sudoku from only inputs and outputs [Amos and Kolter, 2017].

<sup>14</sup> The number of solution variables scales with the number of primal variables, but also the number of constraints. This potentially makes applying the IFT to optimization problems with exponentially many constraints difficult.

<sup>15</sup> For unrolled AD, we could not fully solve the optimization problem and use the truncated trace. For the IFT, vector-inverse Hessian vector product approximations other than the Neumann series include conjugate gradient, identity approximations, and other factorizations such as K-fac [Martens and Grosse, 2015].

<sup>16</sup> The argmax problem is given by maximize  $z^\top \theta$ , subject to  $z^\top \mathbf{1}_n = 1$  and  $z \succeq 0$ , otherwise known as the support function of the constraint set.

OptNet applies the IFT to quadratic programs (QPs) in particular. As the simplest nonlinear optimization problem, QPs strike a balance between expressivity and computational tractability [Frank and Wolfe, 1956]. The methodology remains the same as the softmax problem: Given a QP and a solution, use the KKT conditions to produce a system of equations then apply the IFT to compute the derivative of the solution wrt the parameters of the objective and constraints.

Quadratic programs take the following form:

$$\begin{aligned} & \text{minimize} && \frac{1}{2}z^\top Qz + q^\top z \\ & \text{subject to} && Gz \leq h \\ & && Az = b, \end{aligned} \tag{11}$$

where we optimize over  $z \in \mathbb{R}^n$  and the parameters are  $\theta = \{Q, q, A, b, G, h\}$ , with  $Q \in \mathbb{R}^{n \times n} \succeq 0$ ,  $q \in \mathbb{R}^n$ ,  $A \in \mathbb{R}^{m \times n}$ ,  $b \in \mathbb{R}^m$ ,  $G \in \mathbb{R}^{p \times n}$ ,  $h \in \mathbb{R}^p$ . Compared to the softmax problem in Eqn. 2, the main difference is the learnable parameters in the constraints. As the application of the IFT is very similar, we will not cover it in as much detail as the softmax problem.

Similar to the softmax problem, we will first assume we already have a solution. Unlike in the softmax problem, we assume the the solution to the QP is provided by a solver, as there may not be a closed-form solution.<sup>17</sup> In this case, the execution trace of the solver may be too large to store in memory, necessitating the use of the IFT.

## Step 2: The KKT conditions determine the system of equations

The Lagrangian of the QP is given by

$$\mathcal{L}(\theta, z, u, v) = \frac{1}{2}z^\top Qz + q^\top z + u^\top (Gz - h) + v^\top (Az - b),$$

introducing dual variables  $u \in \mathbb{R}^p$ ,  $v \in \mathbb{R}^m$ , where we have the solution  $x = (z, u, v)$ .

We use the following subset of the KKT conditions to determine our system of equations: stationarity, primal feasibility (equality), and complementary slackness. This yields, written in full,

$$\begin{aligned} Qz + q + G^\top u + A^\top v &= \mathbf{0}_n \\ \text{diag}(u)(Gz - h) &= \mathbf{0}_p \\ \text{diag}(v)(Az - b) &= \mathbf{0}_m. \end{aligned} \tag{12}$$

## Step 3: Check the conditions of the IFT

As with the softmax problem, the only nontrivial part is checking that the Jacobian matrix of  $\frac{dF(\theta, x)}{dx}$  is not singular at the solution point(s) of interest  $(\theta, x)$ . This computation is more involved for QPs than the softmax problem. See Theorem 1 and its proof in the OptNet paper for more details [Amos and Kolter, 2017].

<sup>17</sup> One contribution of OptNet was the extension of a state-of-the-art interior point solver [Amos and Kolter, 2017], and its adaptation to parallel machines (GPUs) and batch processing. While outside the scope of this report, see the paper by Amos and Kolter [2017] for the details.

**Step 4: Compute  $\frac{dx}{d\theta}$**

We would then compute  $\frac{dx}{d\theta} = \left[ \frac{dF(\theta, x)}{dx} \right]^{-1} \frac{dF(\theta, x)}{d\theta}$  by hand or numerically. Amos and Kolter [2017] also show how to obtain the vector-Jacobian product efficiently using quantities readily available from a QP solver.

## 6 Speedups and Spacedowns

In this section we will revisit the softmax problem and experiment with whether the derivative can be approximated more efficiently in either time or space. First we will review the efficient vector-Jacobian product (VJP), then attempt to apply different approximations in the IFT to see if we can find a sublinear algorithm.

Recall softmax,

$$\text{softmax}(\theta)_i = \frac{\exp(\theta_i)}{\sum_j \exp(\theta_j)},$$

and the softmax problem from Eqn. 2,

$$\begin{aligned} \text{maximize} \quad & z^\top \theta + H(z) \\ \text{subject to} \quad & z^\top \mathbf{1} = 1 \\ & z_i \geq 0, \forall i. \end{aligned} \tag{13}$$

## 7 Gradient Estimation in Structured Attention

We now apply the IFT to structured latent variable models.

Similar to the softmax case, we can compute both the solution and gradient exactly without using too much memory. In the case of hidden Markov models (HMMs), we can perform inference in time linear in the length of a sequence and quadratic in the number of states and space linear in length and the number of states.

## 8 Variational Inference

Variational inference has found success in recent applications to generative models, in particular by allowing practitioners to depart from conjugate models and extend emission models with expressive neural network components. The main insight that led to this development is that inference can be amortized through the use of an inference network. One approach to variational inference, stochastic variational inference (SVI), introduces local, independent variational parameters for every instance of hidden variable. While flexible, the storage of all variational parameters is expensive, and the optimization of each parameter independently slow []. Amortized variational inference (AVI) solves that by instead using a hierarchical process. Variational parameters are produced hierarchically via an inference network, which in turn generates the local variational parameters []. The resulting local parameters may or may not be subsequently optimized.

Failure to further optimize local variational parameters may result in an amortization gap []. Prior work has shown that this gap can be ameliorated by performing a few steps of optimization on the generated local parameters obtained from the inference network, and even by propagating gradients through the optimization process. Optimizing through the inner optimization problem results in semi-amortized variational inference (SAVI) [].

As our main motivating example, we will examine whether we can apply the IFT to SAVI. We will start by formalizing the problem of variational inference for a simple model.

We will start with a model defined by the following generative process, used by Dai et al. [2019] to analyze posterior collapse:

1. Choose a latent code from the prior distribution  $z \sim p(z) = N(0, I)$ .
2. Given the code, choose an observation from the emission distribution  $x | z \sim p_\theta(x | z) = N(\mu_x(z, \theta), \gamma I)$ ,

where  $\mu_x(z, \theta) \equiv \text{MLP}(z, \theta)$  and  $\gamma > 0$  is a hyperparameter. This yields the joint distribution  $p(x, z) = p(x | z)p(z)$ .

Since the latent code  $z$  is unobserved, training this model would require optimizing the evidence  $p(x) = \int p(x, z)$ . However, due to the MLP parameterized  $\mu_x$ , the integral is intractable. Variational inference performs approximate inference by introducing variational distribution  $q_\phi(z | x)$  and maximizing the following lower bound on  $\log p(x)$ :

$$\log p(x) - D_{\text{KL}}[q(z | x) || p(z | x)] = \mathbb{E}_{q_\phi(z | x)} \left[ \log \frac{p_\theta(x, z)}{q_\phi(z | x)} \right] = L(\theta, \phi). \quad (14)$$

(Write out objective in full.)

While SVI introduces local parameters for each instance of  $z$ , and AVI uses a single  $q(z | x)$  for all instances, we will follow the approach of SAVI. We will perform inference as follows: For each instance  $x$ , produce local variational parameter  $z^{(0)} = g(x; \phi)$ . Obtain  $z^*$  by solving  $\mathcal{L}(\theta, z^{(0)}) = 2$ , with (local) optima  $\ell^*$ . Take gradients through the whole procedure, i.e. compute  $\frac{\partial \ell^*}{\partial \phi} = \frac{\partial \ell^*}{\partial z^*} \frac{\partial z^*}{\partial z^{(0)}} \frac{\partial z^{(0)}}{\partial \phi}$ . The main difficulty lies in computing  $\frac{\partial z^*}{\partial z^{(0)}}$ . (Highlight challenge)

In order to avoid the memory costs of storing all intermediate computation performed in a solver, we will instead apply the IFT. In order to apply the IFT, we must satisfy the three conditions. First, we must have a solution point to a system of equations,  $F(x_0, z_0) = 0$ . In this setting, we will use the KKT conditions of the optimization problem to define  $F$ .

## 9 Limitations

## References

A. Agrawal, B. Amos, S. T. Barratt, S. P. Boyd, S. Diamond, and J. Z. Kolter. Differentiable convex optimization layers. *CoRR*, abs/1910.12430, 2019. URL <http://arxiv.org/abs/1910.12430>.

- 331 B. Amos and J. Z. Kolter. Optnet: Differentiable optimization as a layer in neural networks. *CoRR*,  
332 abs/1703.00443, 2017. URL <http://arxiv.org/abs/1703.00443>.
- 333 Q. Bertrand, Q. Klopfenstein, M. Blondel, S. Vaiter, A. Gramfort, and J. Salmon. Implicit differen-  
334 tiation of lasso-type models for hyperparameter optimization, 2020.
- 335 S. Boyd and L. Vandenberghe. *Convex Optimization*. Cambridge University Press, March 2004.  
336 ISBN 0521833787. URL [http://www.amazon.com/exec/obidos/redirect?tag=](http://www.amazon.com/exec/obidos/redirect?tag=citeulike-20&path=ASIN/0521833787)  
337 [citeulike-20&path=ASIN/0521833787](http://www.amazon.com/exec/obidos/redirect?tag=citeulike-20&path=ASIN/0521833787).
- 338 B. Dai, Z. Wang, and D. P. Wipf. The usual suspects? reassessing blame for VAE posterior collapse.  
339 *CoRR*, abs/1912.10702, 2019. URL <http://arxiv.org/abs/1912.10702>.
- 340 C. Finn, P. Abbeel, and S. Levine. Model-agnostic meta-learning for fast adaptation of deep net-  
341 works. *CoRR*, abs/1703.03400, 2017. URL <http://arxiv.org/abs/1703.03400>.
- 342 M. Frank and P. Wolfe. An algorithm for quadratic programming. *Naval Research Logis-*  
343 *tics Quarterly*, 3(1-2):95–110, 1956. URL [https://EconPapers.repec.org/RePEc:](https://EconPapers.repec.org/RePEc:wly:navlog:v:3:y:1956:i:1-2:p:95-110)  
344 [wly:navlog:v:3:y:1956:i:1-2:p:95-110](https://EconPapers.repec.org/RePEc:wly:navlog:v:3:y:1956:i:1-2:p:95-110).
- 345 B. Gao and L. Pavel. On the properties of the softmax function with application in game theory and  
346 reinforcement learning, 2018.
- 347 A. N. Gomez, M. Ren, R. Urtasun, and R. B. Grosse. The reversible residual network: Backpropa-  
348 gation without storing activations. *CoRR*, abs/1707.04585, 2017. URL [http://arxiv.org/](http://arxiv.org/abs/1707.04585)  
349 [abs/1707.04585](http://arxiv.org/abs/1707.04585).
- 350 S. Gould, R. Hartley, and D. Campbell. Deep declarative networks: A new hope. *CoRR*,  
351 abs/1909.04866, 2019. URL <http://arxiv.org/abs/1909.04866>.
- 352 A. Griewank and A. Walther. *Evaluating Derivatives: Principles and Techniques of Algorithmic*  
353 *Differentiation*. SIAM, Philadelphia, 2nd edition, 2008.
- 354 W. Karush. *Minima of functions of several variables with inequalities as side conditions*. PhD thesis,  
355 Thesis (S.M.)—University of Chicago, Department of Mathematics, December 1939., 1939.
- 356 Y. Kim, S. Wiseman, A. C. Miller, D. Sontag, and A. M. Rush. Semi-amortized variational autoen-  
357 coders, 2018.
- 358 S. Krantz and H. Parks. The implicit function theorem : History, theory, and applications / s.g.  
359 krantz, h.r. parks. 01 2003. doi: 10.1007/978-1-4612-0059-8.
- 360 H. W. Kuhn and A. W. Tucker. Nonlinear programming. In *Proceedings of the Second Berkeley*  
361 *Symposium on Mathematical Statistics and Probability, 1950*, pages 481–492, Berkeley and Los  
362 Angeles, 1951. University of California Press.
- 363 J. Lorraine, P. Vicol, and D. Duvenaud. Optimizing millions of hyperparameters by implicit differ-  
364 entiation. *CoRR*, abs/1911.02590, 2019a. URL <http://arxiv.org/abs/1911.02590>.

- 365 J. Lorraine, P. Vicol, and D. Duvenaud. Optimizing millions of hyperparameters by implicit differ-  
 366 entiation. *CoRR*, abs/1911.02590, 2019b. URL <http://arxiv.org/abs/1911.02590>.
- 367 D. Maclaurin, D. Duvenaud, and R. P. Adams. Gradient-based hyperparameter optimization through  
 368 reversible learning, 2015.
- 369 J. Martens and R. B. Grosse. Optimizing neural networks with kronecker-factored approximate  
 370 curvature. *CoRR*, abs/1503.05671, 2015. URL <http://arxiv.org/abs/1503.05671>.
- 371 A. F. T. Martins and R. F. Astudillo. From softmax to sparsemax: A sparse model of attention and  
 372 multi-label classification. *CoRR*, abs/1602.02068, 2016. URL [http://arxiv.org/abs/](http://arxiv.org/abs/1602.02068)  
 373 [1602.02068](http://arxiv.org/abs/1602.02068).
- 374 A. Rajeswaran, C. Finn, S. M. Kakade, and S. Levine. Meta-learning with implicit gradients. *CoRR*,  
 375 abs/1909.04630, 2019. URL <http://arxiv.org/abs/1909.04630>.
- 376 S. Venkataraman and B. Amos. Neural fixed-point acceleration for convex optimization. *CoRR*,  
 377 abs/2107.10254, 2021. URL <https://arxiv.org/abs/2107.10254>.
- 378 M. Wainwright and M. Jordan. Graphical models, exponential families, and variational inference.  
 379 *Foundations and Trends in Machine Learning*, 1:1–305, 01 2008. doi: 10.1561/22000000001.

## 380 A Properties of Exponential Family Distributions

381 Exponential family distributions are widely used to parameterize graphical models [Wainwright  
 382 and Jordan, 2008]. The classification of a distribution as exponential family simply says that we  
 383 can write their joint distributions with the particular factorization given above, which can then be  
 384 used to derive useful properties common to all exponential family distributions. These properties,  
 385 made clearer through the exponential family abstraction, are then integral for performing inference.  
 386 In particular, gradient estimation can be written as mapping from canonical to mean parameters.

### 387 A.1 Canonical Parameters

Exponential family distributions take the form

$$p(x; \theta) = h(x) \exp(\theta^\top \phi(x) - A(\theta)), \quad (15)$$

388 where  $h : \mathcal{X} \rightarrow \mathbb{R}_+$  is the base measure,  $\theta \in \mathbb{R}^d$  the canonical parameters,  $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$   
 389 the sufficient statistics<sup>18</sup>, and  $A : \mathbb{R}^d \rightarrow \mathbb{R}$  the log partition function. Intuitively, the canonical  
 390 parameters  $\theta$  score configurations  $x \in \mathcal{X}$ , with representation  $\phi(x)$ . The log partition function  
 391  $A(\theta) = \log \int_{\mathcal{X}} h(x) \exp(\theta^\top \phi(x))$  is used to ensure normalization of the probability density, i.e.  
 392  $\int_{\mathcal{X}} p(x) = 1$ . The base measure  $h : \mathcal{X} \rightarrow \mathbb{R}$  is also used to ensure proper normalization, but is only

---

<sup>18</sup>Potential functions usually refer to particular pairs of canonical parameter and sufficient statistic products  $\theta_\alpha \phi_\alpha(x)$  as a function of  $x$ .



a function of  $x$ . It can also be used to enforce hard constraints in models by giving 0 mass to some configurations  $x$ .

We give three examples of exponential family distributions: softmax, hidden Markov models (HMMs) and linear chain conditional random fields (CRFs).

**Softmax** Softmax gives a discrete distribution over  $n$  items, where  $x \in [1, \dots, n]$ . The probability mass function is then  $p(x; \theta) = \exp(\theta^\top x - A(\theta))$ , where  $A(\theta) = \log \sum_i \exp(\theta_i) = \log \exp(\theta)^\top \mathbf{1}_n$ . Thus, the sufficient statistics are  $\phi(x)_i = 1(x_i)$  for each  $i \in [n]$  and the base measure constant  $h(x) = 1$ .

**Hidden Markov Models (HMMs)** Hidden Markov models (HMMs) give a discrete distribution over sequences  $x = (x_1, \dots, x_T)$ , where  $x_t \in [X]$ , and latent states  $z = (z_1, \dots, z_T)$ , where  $z_t \in [Z]$ . HMMs have the joint distribution  $p(x, z) = \prod_t p(x_t | z_t) p(z_t | z_{t-1})$ , with distinct start state  $z_0 = \epsilon$ . We assume each condition distribution is parameterized using softmax, and the HMM has canonical parameters  $\theta = (O, S)$  for the emission and transitions respectively. The emissions are locally normalized, given by  $p(x_t | z_t) = O_{z_t, x_t}$  where  $O \in [0, 1]^{ZX}$ ,  $1_{z_t, x_t} \in \mathbb{R}^{ZX}$  is a one-hot vector. Local normalization implies  $\sum_x O_{z, x} = 1$  for a given  $z$ . The transitions are also locally normalized, given by  $p(z_t | z_{t-1}) = S_{z_{t-1}, z_t}$ , where  $S \in [0, 1]^{Z^2}$ ,  $1_{z_{t-1}, z_t} \in \{0, 1\}^{Z^2}$  is a one-hot vector. Both  $O$  and  $S$  are vector representations of the observation and state transition matrices respectively. The sufficient statistics are  $\{1(x_t, z_t)\}_{x_t, z_t}$  and  $\{1(z_{t-1}), 1(z_t)\}_{z_{t-1}, z_t}$  for all  $t, x_t, z_{t-1}, z_t$ . The base measure is constant  $h(x, z) = 1$ .

**Linear Chain Conditional Random Fields (CRFs)** The linear chain CRF has states  $x = (x_1, \dots, x_T)$ , where  $x_t \in [Z]$ , and joint distribution

$$p(x) = \frac{\prod_t f(z_t; \theta) g(z_{t-1}, z_t; \theta)}{\exp(A(\theta))} = \exp(\theta^\top \phi(x) - A(\theta)),$$

where  $f : \mathcal{X} \rightarrow \mathbb{R}_+$ ,  $g : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_+$  are unary and binary potentials respectively. These potentials are the most common representation of CRFs. In the exponential family representation, the canonical parameters  $\theta \in \mathbb{R}^{(Z+Z^2)T}$  are used to score unary and binary terms, given by the sufficient statistics  $\phi(x)_{t,1,i} = 1(z_t = i)$  and  $\phi(x)_{t,2,ij} = 1(z_{t-1} = i, z_t = j)$  for  $t, i, j$ . The base measure is constant  $h(x) = 1$ .

## A.2 Mean Parameters

The mean parameters of exponential family distributions are the expected sufficient statistics:  $\mu = \mathbb{E}_{p(x)} [\phi(x)]$ . The space of all possible marginals, obtained by varying the distribution  $p(x)$ , is called the mean parameter space

$$\mathcal{M} = \{\mu \mid \exists p \text{ s.t. } \mathbb{E}_p [\phi(x)] = \mu\}.$$

For the particular case of discrete graphical models, which is of particular interest, we call  $\mathcal{M}$  the marginal polytope. In particular, for discrete models, we have

$$\mathcal{M} = \left\{ \mu \mid \sum_x p(x) \phi(x) = \mu, \sum_x p(x) = 1 \right\},$$

expressed as a finitely generated polyhedra. By the Minkowski-Weyl theorem, we can equivalently express the marginal polytope as a constrained polyhedra, i.e. the convex hull of the sufficient statistics:

$$\mathcal{M} = \text{conv}(\phi(x), x \in \mathcal{X}).$$

418 This representation is more useful in variational optimization problems, as it translates directly to a  
419 set of constraints.

### 420 **A.3 Forward and Backward Mappings from Canonical to Mean Parameters**

421 The forward mapping  $\theta \mapsto \mu$  is marginal inference in CRFs, while the backward mapping  $\mu \mapsto \theta$   
422 corresponds to learning. The log partition function  $A(\theta)$  is key to both of these mappings:

423 **Forward**  $\nabla A(\theta) = \mu$

**Backward**  $\theta \in \nabla A^*(\mu)$

425 where  $A^*(\mu)$  denotes the conjugate dual function of  $A$ . When the mean parameters are in the  
426 interior of the mean parameter space (henceforth referred to as the marginal polytope, as we focus  
427 on discrete graphical models)  $\mu \in \mathcal{M}^\circ$ , the conjugate dual function  $A^*(\mu) = -H(p_{\theta(\mu)})$  **More**  
428 **precise definition following Wainwright and Jordan [2008] Thm. 3.4.**

## 429 **B Graphical Models**

430 Neural ODEs use reversibility.